



(a) Interleaved duet representation



each stream keeps a private history;
stream identity via per-stream Q/K/V,
not position


 melody frame
 chord frame
 SOS
 diffusion query slot (noise level k)
 empty frame (sparse stream)

(b) Training: one block-diffusion step



slots share frame T's rotary phase $\rightarrow$ training phase $\equiv$ decoding phase at every t (no padding)
loss = AR (all clean rows) + denoise (slots, random T and k per sample)

(c) Decoding: seeded refinement per frame



within a frame, exposure bias is closed by construction:
the slots condition on their own estimates (as in training
via self-conditioning), not on ground truth